AI PROJECT Dino Game Playing Using Reinforcement Learning

Presented by:-

- Arzoo Jangra 6
- Dhirender Kumar 10
- Vidhi Khare 54

Dino Run or T-Rex run is an endless runner game in Chrome browser which is available to play when you're offline aka 'the game you don't usually like to see'.



Technologies Used

Python

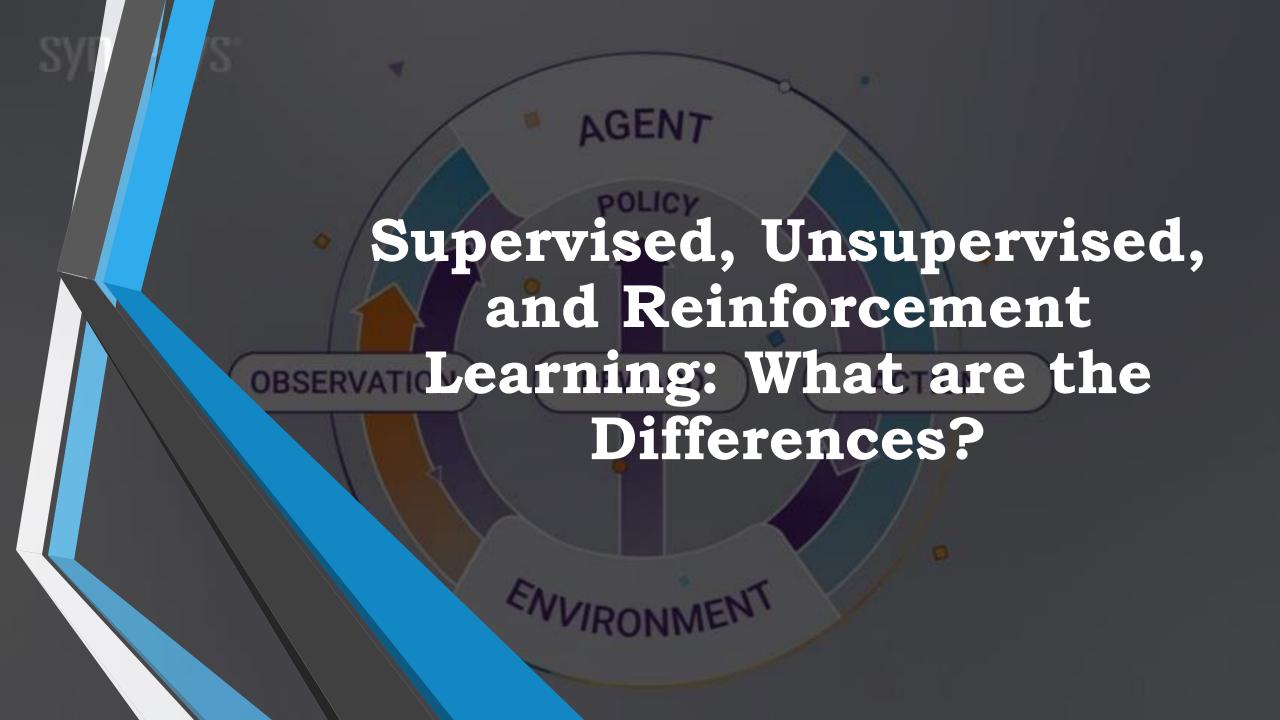
- Python Selenium Package
- Numpy, Matplotlib, Pandas, Seaborn

Reinforcement learning

• Q-Learning

Git/Github

- O **NumPy** can be used to perform a wide variety of mathematical operations on arrays.
 - Matplotlib is a cross-platform, data visualization and graphical plotting library for Python.
 - O Pandas is a Python library for data analysis.
 - **Seaborn** is a library in Python predominantly used for making statistical graphics.
 - Selenium, a popular browser automation tool, was used to send actions to the browser and get different game parameters like current score.



Difference #1: Static Vs. Dynamic

The goal of supervised and unsupervised learning is to search for and learn about patterns in training data, which is quite static. RL, on the other hand, is about developing a policy that tells an agent which action to choose at each step — making it more dynamic.

Difference #2: No Explicit Right Answer

In supervised learning, the right answer is given by the training data. In Reinforcement Learning, the right answer is not explicitly given: instead, the agent needs to learn by trial and error. The only reference is the reward it gets after taking an action, which tells the agent when it is making progress or when it has failed.

Difference #3: RL Requires Exploration

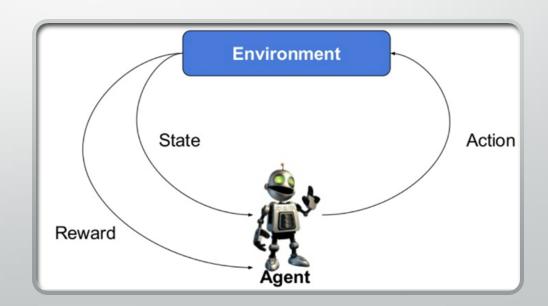
A Reinforcement Learning agent needs to find the right balance between exploring the environment, looking for new ways to get rewards, and exploiting the reward sources it has already discovered. In contrast, supervised and unsupervised learning systems take the answer directly from training data without having to explore other answers.

Difference #4: RL is a MultipleDecision Process

Reinforcement Learning is a multiple-decision process: it forms a decision-making chain through the time required to finish a specific job. Conversely, supervised learning is a single-decision process: one instance, one prediction.

Reinforcement Learning

• Reinforcement Learning is a feedback-based Machine learning technique in which an agent learns to behave in an environment by performing the actions and seeing the results of actions. For each good action, the agent gets positive feedback, and for each bad action, the agent gets negative feedback or penalty.



Elements of Reinforcement Learning

Reinforcement Learning has **four** essential elements:

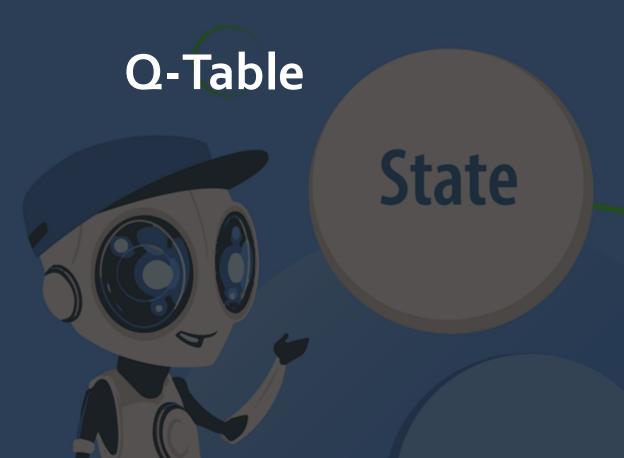
- **Agent:** The program you train, with the aim of doing a job you specify.
- **Environment:** The world, real or virtual, in which the agent performs actions.
- **Action:** A move made by the agent, which causes a status change in the environment.
- **Rewards:** The evaluation of an action, which can be positive or negative.

Q-Learning

- Q-learning is a popular **values- based** reinforcement learning algorithm based on the **Bellman equation**.
- The main objective of Q-learning is to learn the policy which can inform the agent that what actions should be taken for maximizing the reward under what circumstances.
- It is an **off-policy RL** that attempts to find the best action to take at a current state.
- The 'Q' in Q-learning stands for quality. Quality here represents how useful a given action is in gaining some future reward.

- **Agent:** Dino
- **Environment:** Web Browser
- **Action:** Jump & Duck
- **Rewards:** 15 for not getting out
- **Penalty:** -1000 for getting out

Elements
of Reinforcement
Learning
According to our
project

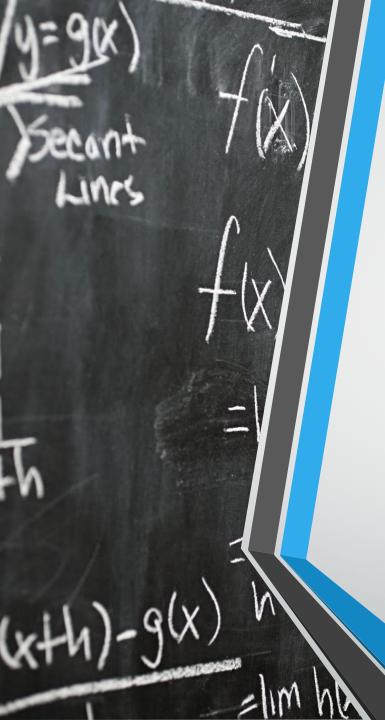


Q Table

Q-Table is the data structure used to calculate the maximum expected future rewards for action at each state. Basically, this table will guide us to the best action at each state.

<u>State</u>							<u>Action</u>		
Dino Information		Obstacle Information		Current speed of game	<u>Do Nothing</u>	<u>Jump</u>	<u>Duck</u>		
	X_Coordinate	Y_Coordinate	X Coordinate	Y_Coordinate	<u>Width</u>				
	45	36	80	33	50	8	3	2.5	-1
	50	34	56	45	65	8	1	3	-4
	55	39	67	34	70	8	2	5.1	-2

Q- Table According to our project



Maths behind Q-Learning

The goal of the agent in Q-learning is to maximize the value of Q-function [Q(s, a)], which uses the Bellman equation and takes two inputs state(s) and action(a).

Notations:

s = A particular state

a = Action

s' = State to which user goes from s

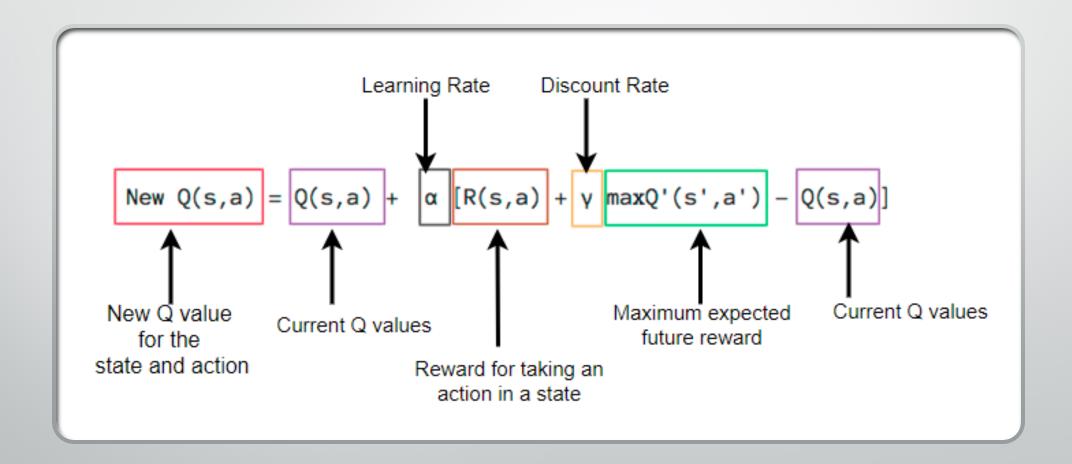
 Υ = Discount factor

 α = Learning rate

R(s, a) = Reward function which takes a state s and action a and outputs a reward value.

Q'(s', a') = Expected future reward.

Q-function



Exploration vs Exploitation

Exploration vs Exploitation problem arises when our model tends to stick to same actions while learning, in our case the model might learn that jumping gives better reward rather than doing nothing and in turn apply an always jump policy. However, we would like our model to try out random actions while learning which can give better reward. We introduce ε , which decides the randomness of actions. We gradually decay its value to reduce the randomness as we progress and then exploit rewarding actions.

Credit Assignment Problem

Credit Assignment problem can confuse the model to judge which past action was responsible for current reward. Dino cannot jump again while mid-air and might crash into a cactus, however, our model might have predicted a jump. So the negative reward was in fact a result of previously taken wrong jump and not the current action. We introduce Discount Factor y, which decides how far into the future our model looks while taking an action. Thus, y solves the credit assignment problem indirectly. In our case the model learned that stray jumps will inhibit it's ability to jump in the future when we set $\gamma=0.99$

Procedure

- ✓ Game Initialization
- ✓ Initialize the Variables and Q-Table
- ✓ Start game playing
- ✓ Get current state of dino
- ✓ Check for the current state in the Q-Table
 - ✓ If current state is in Q-Table get the Q-Values for the given state
 - ✓ Else add this new state to the Q-Table and randomly generate its Q-Values
- ✓ Based on the Q-Values take the action with highest Q-Value
- ✓ Based on the result of the action update the Q-Values for the current (state, action).

```
67 det startGame():
       driver.find element(By.TAG NAME, "body").send keys(Keys.SPACE)
69
70
  def initialize():
72
       s = Service('/home/the-techie/Desktop/chromedriver')
73
74
       options = webdriver.ChromeOptions()
75
       options.add argument("--disable-web-security")
       driver = webdriver.Chrome(service=s, options=options)
76
77
       driver.maximize window()
78
79
       try:
           driver.get("chrome://dino/")
81
       except:
82
           pass
83
84
       return driver
85
86
  def addNewState(state):
       q_table[state] = np.round(np.random.rand(3), decimals = 2)
88
```

Game Initialization

Initialize the Variables and Q-Table

```
import numpy as np
2 from selenium import webdriver
3 from selenium.webdriver.chrome.service import Service
4 from selenium.webdriver.common.keys import Keys
5 from selenium.webdriver.common.by import By
6 from selenium.webdriver.support.ui import WebDriverWait
8 scores = []
9 q table = dict()
10 \text{ reward} = 15
11 out = -1000
12 learning rate = 0.1
  discount = 0.99
14 \text{ episodes} = 20000
15 driver = None
16
```

Start game playing

Get current state of Dino

```
def getState():
       x pos = 0
31
       y pos = 0
       obs pos = []
32
33
       curr speed = float(driver.execute script("return Runner.instance .currentSpeed"))
35
       if round(curr speed % 1, 1) >= 0.5:
           curr speed = int(curr speed) + 0.5
37
       else:
           curr speed = float(int(curr speed))
       x pos = driver.execute script("return Runner.instance .tRex.xPos")
       y pos = driver.execute script("return Runner.instance .tRex.yPos")
41
       isJumping = driver.execute script("return Runner.instance .tRex.jumping")
42
43
       num obstacles = int(driver.execute script("return Runner.instance .horizon.obstacles.length"))
44
       if num obstacles > 0:
45
           a = min(150, driver.execute script("return Runner.instance .horizon.obstacles[0].xPos")//3)
47
           b = driver.execute script("return Runner.instance .horizon.obstacles[0].yPos")//2
           width = driver.execute script("return Runner.instance .horizon.obstacles[0].width")
           obs pos = [a, b, width]
51
52
        res = tuple([x pos, y pos, isJumping, round(curr speed, 1)] + obs pos)
54
        return res
```

```
def addNewState(state):
        q table[state] = np.round(np.random.rand(3), decimals = 2)
86
87
   def getQvalues(curr state):
89
        try:
            p = q_table[curr_state]
90
91
92
        except:
            addNewState(curr state)
            p = q table[curr state]
94
95
96
        return p
```

Get the Q-Value for the current state

Based on the Q-Values take the action with highest Q-Value

```
17 # 0: Nothing
18 # 1: JUMP
19 # 2: DUCK
20
21 def performAction(choice):
22    if choice == 1:
23         driver.find_element(By.TAG_NAME, "body").send_keys(Keys.ARROW_UP)
24    elif choice == 2:
25         driver.find_element(By.TAG_NAME, "body").send_keys(Keys.ARROW_DOWN)
26
27
```

```
p = getQvalues(curr_state)

129

130

131

132

performAction(action)

133
```

Based on the result of the action update the Q-Values for the current (state, action)

```
# updating *****************
                isPlaying = driver.execute script("return Runner.instance .playing")
138
                if isPlaying:
                    reward = 15
                else:
                    reward = -1000
                current qsa = p[action]
                next state = getState()
                next state qs = None
                try:
                    next state qs = q table[next state]
                except:
                    addNewState(next state)
                    next state qs = q table[next state]
                max future q = np.max(next state qs)
                new q = (1 - learning rate ) * current qsa + learning rate * (reward + discount * max future q)
                p[action] = new q
                q table[curr state] = p
                curr state = next state
```

Week 2 updates (17-01-2022 to 23-01-2022)

State Current speed of game

6.5

8

6

Dino Information

Y-Coordinate

82

93

93

70

Is Jumping

<u>(T/F)</u>

F

F

F

X-Coordinate

0

23

23

23

Action

<u>Jump</u>

0.36

2.36

0.21

0.1

<u>Duck</u>

0.39

0.87

0.27

0.2

Do Nothing

1.96

0.62

2.19

0.51

Obstacle Information

Y-Coordinate

52

52

52

Width

51

51

51

X-Coordinate

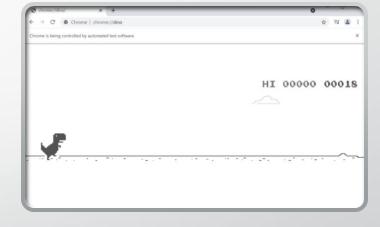
150

78

14

Explanation

- We are performing a certain number of episodes in one run of the program.
- At start of the game, we get the initial state of the Dino and the Obstacles(if present). Since in the image it can be seen that there is no obstacle present so our program only returns the state of the Dino (i.e. it's x-coordinate, y-coordinate, is it jumping or not, current speed of the game).





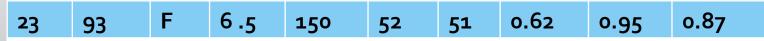
• After getting the state, we retrieve the current score from the browser and while the score is less than 35, our Dino will do nothing.

- If score>35 and our Dino is still playing then we try to get the Q-Value for the current state using **getQvalues()** method.
- If current state of our Dino is not present in the Q-Table then we will add that state into the Q-Table by generating random Q-value for that state and generated Q-values would be returned by our method.
 - In our case it returns 0.45(for Nothing), 0.36(for Jump), 0.39(for Duck)
 - O Current State -> o 82 T 6 0.45 0.36 0.39
- If the current state is already present in the Q-Table then the **getQvalues()** would simply retrieve the Q-values for the current state from the Q-Table and return it.
- Further the Dino would perform action based on the action with highest Q-values.
 - O In our case the max Q-value is 0.45 (Do Nothing Index 0)

- O After taking the action we check if our Dino is still playing or not.
 - O If alive we give Dino reward of 15
 - Else a penalty of –1000

If alive (Reward = 15)

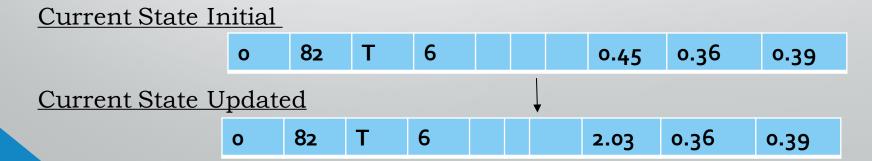
- We update the Q-value of the particular action taken for the current state to incur the effect of reward.
- To calculate the updated Q-value we retrieve the state in which Dino is after performing the action.
- We retrieve the next state, Q-values for the next state(next_state_qs) and also the maximum of the Q-values of this next state(max_future_q).
 - O Next_state



- \bigcirc Next_state_qs = [0.62, 0.95, 0.87]
- Max_future_q = 0.95 for Jump

Calculation to Update Q-Value, If Dino is Still playing

- Now we calculate the updated Q-value using the Q-Function
- new_q = [(1 learning_rate) * current_qsa] +[learning_rate * (reward + discount * max_future_q)]
 - new_q = [(1 0.1) * 0.45] + [0.1 * (15 + 0.99 * 0.95)]
 - new_q = [0.99*1.27] + [0.1*15.9405]
 - new_q = 0.4455 + 1.59405
 - new_q = 2.03
 - Now we update this calculated new_q in the q_table for the current state and for the action that was taken (I.e. Do Nothing)

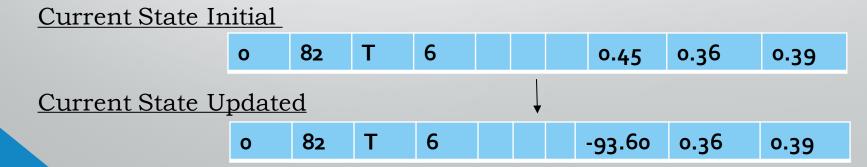


Else (Reward = -1000)

- Here the reward represents the penalty given to the Dino.
- Penalty is given if the Dino gets out in the game.
- We update the Q-value of the particular action taken for the current state to incur the effect of penalty.
- To calculate the updated Q-value we retrieve the state in which Dino is after performing the action.
- We retrieve the next state, Q-values for the next state(next_state_qs) and also the maximum of the Q-values of this next state(max_future_q).
 - Next_state
 23 93 F 6.5 150 52 51 0.62 0.95 0.87
 - \bigcirc Next_state_qs = [0.62, 0.95, 0.87]
 - Max_future_q = 0.95 for Jump

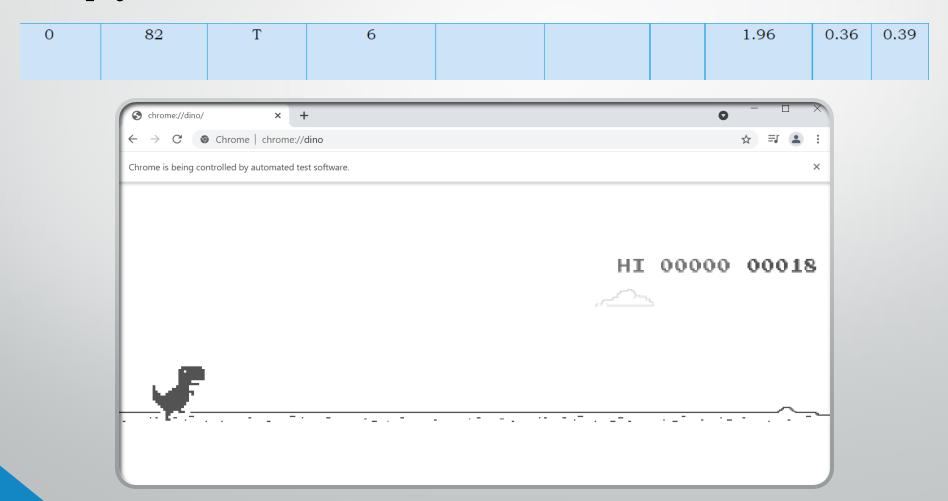
Calculation to Update Q-Value, If Dino is Still playing

- Now we calculate the updated Q-value using the Q-Function
- new_q = [(1 learning_rate) * current_qsa] +[learning_rate * (reward + discount * max_future_q)]
 - new_q = [(1 0.1) * 0.45] + [0.1 * (-1000 + 0.99* 0.95)]
 - new_q = [0.99*1.27] + [0.1*(-940.5)]
 - new_q = 0.4455 94.05
 - $new_q = -93.6045$
 - Now we update this calculate new_q in the q_table for the current state and for the action that was taken (I.e. Do Nothing)



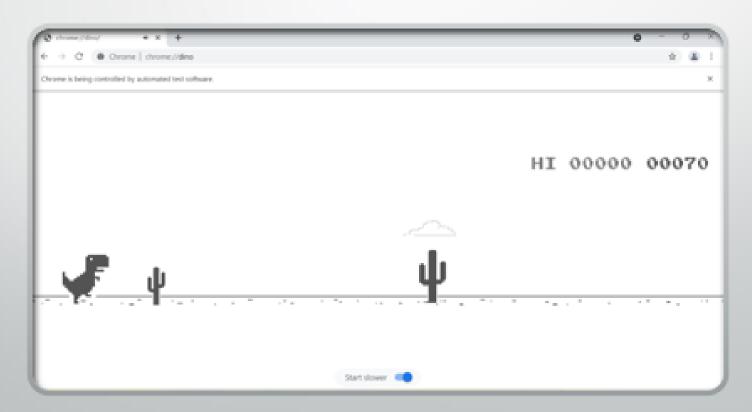
- Further if Dino is alive we would make next state as our current state to work over the next state that we are in after taking the action.
- We would go on performing the method listed in previous slides for all the upcoming states and fill up our Q-Table simultaneously to train our Dino.
- If Dino gets out, a different episodes starts with the same process listed but with a filled Q-Table which it uses to train Dino.
- After each episode we are storing the score that is generated and printing the same.

• As we can see in the image that there is no obstacle on the screen in the very beginning of the game that's why the obstacle info part of the first row of our Q-table is empty.



• Now in the picture below we can see the obstacle arriving towards our Dino and hence the second row of our Q-table provides the information about the obstacle as well which helps it to perform action accordingly.





Training in batches

- Previously, we had to train the model in one go.
- Now, with the use of json file, we are able to train our model in batches and on multiple systems.

Json File

This file contains our Q-table for all the sessions which helps us to start a next session with this filled Q-table instead of starting with an empty Q-table which also helps us to train our model in different batches on different systems.

```
{"0 82 True 6 ": [1.96, 0.36, 0.39], "23 93 False 6.5 150 52 51": [0.62, 2.36, 0.87], "23 93 False 6.5 78 52 51": [2.19, 0.21, 0.27], "23 93 False 6.5 74 52 51": [2.14, 0.46, 0.39], "23 93 False 6.5 70 52 51": [0.69, 0.54, 2.27], "23 93 False 6.5 63 52 51": [0.04, -99.33, 0.3], "23 64 True 6.5 -2 52 51": [0.58, 0.71, 0.64], "23 93 False 6 ": [13.75, 0.32, 0.53], "23 93 False 6 92 45 25": [0.18, -99.27, 1.84], "23 80 True 6 16 45 25": [0.97, 0.89, 0.29], "23 93 False 6 90 45 25": [0.71, 0.33, 2.41], "23 93 False 6 70 45 25": [0.36, 3.52, 0.06], "23 93 False 6 99 52 51": [0.52, -99.04, 0.88], "23 93 False 6 18 52 51": [-99.16, 0.11, -99.18], "23 93 False 6 89 52 [0.05, -99.12, 0.47], "23 70 True 6 14 52 51": [0.51, 0.1, 0.2], "23 93 False 6 90 52 17": [3.77, 0.85, 0.2], "23 93 False 6 78 52 17": [0.6, 2.44, 0.66], "23 93 False 6 80 45 25": [0.69, -99.04, 0.74], "23 53 True 6 8 45 25": [0.66, 0.21, 0.44], "23 93 False 6 92 52 51": [0.21, -99.66, 0.12], "23 78 True 6 15 52 51": [0.41, 0.57, 0.52], "23 93 False 6 91 45 25": [0.39, 0.65, 2.23], "23 93 False 6 62 45 25": [2.03, 0.11, 0.14], "23 93 False 6 48 45 25": [2.19, 0.42, 0.5], "23 93 False 6 34 45 25": [0.41, 0.26, -99.36], "23 93 False 6 20 45 25":
```

Json file snippet

Week 3 updates

(14-02-2022 to 20-02-2022)

Training Phase

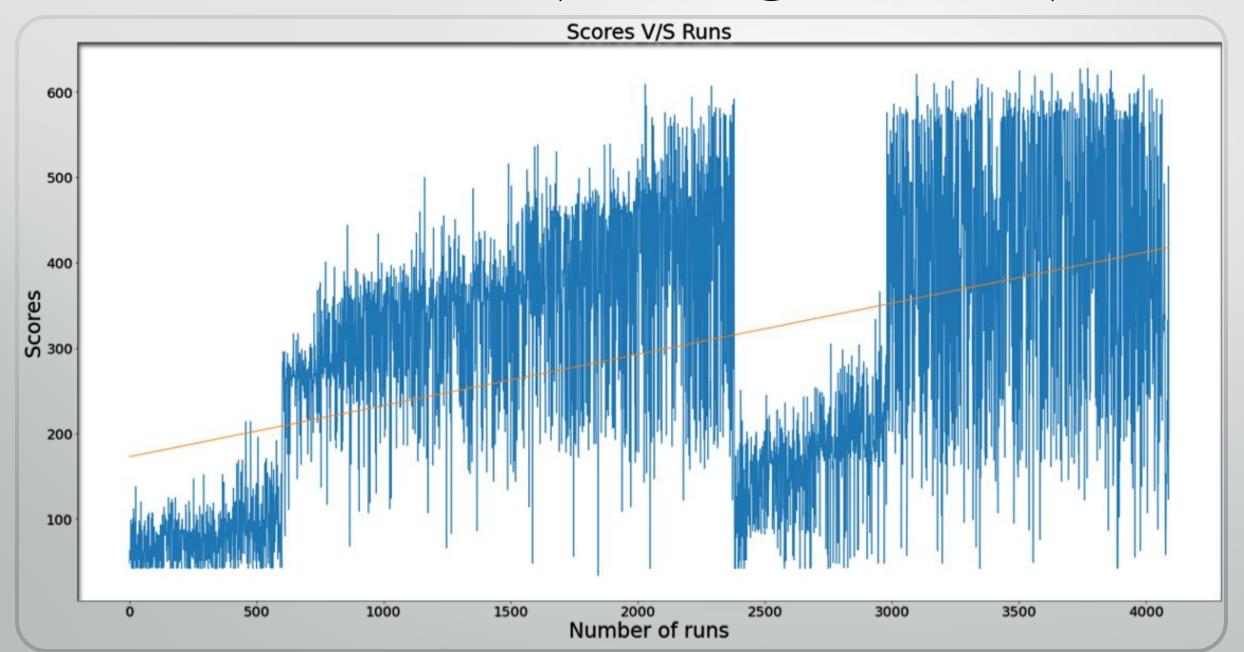
- What are the parameter changes done for training to get the best results?
 - We are changing the learning rate in our program for training in different phases. (0.1 to 0.6)

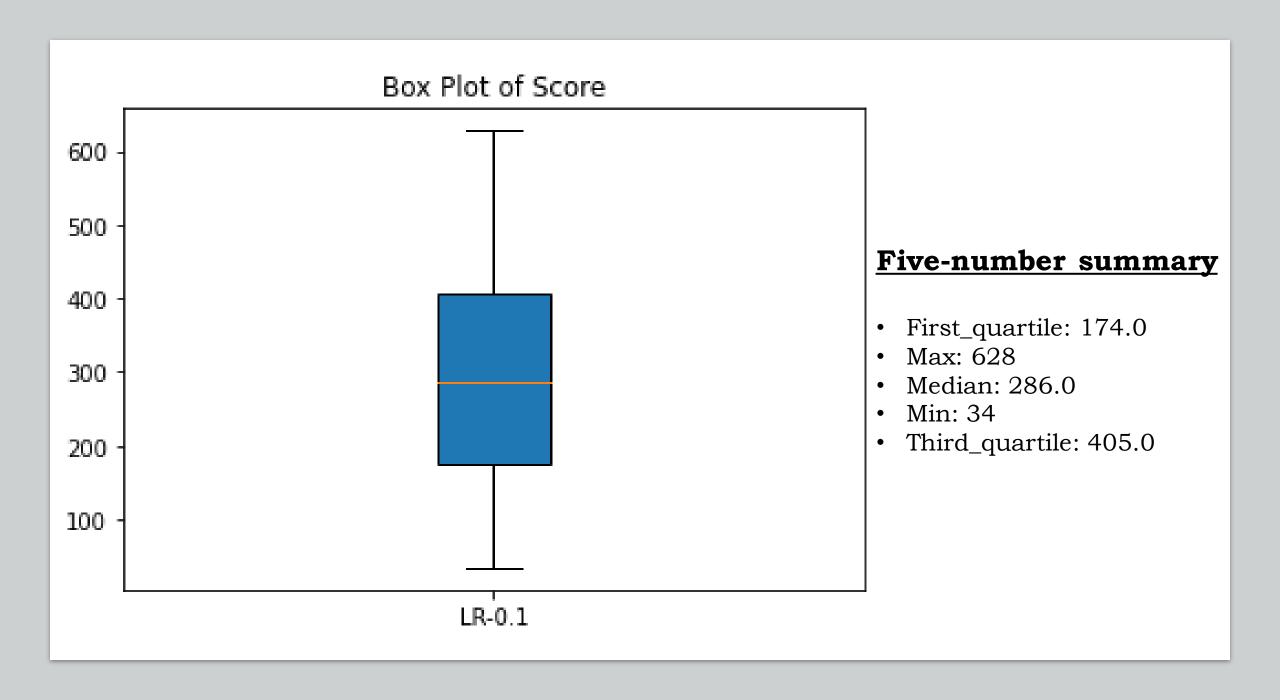
Why we choose different values of learning rate for training?

Since it can be seen that learning rate is making a tradeoff between our current q-value and future q-value, thereby to get the effect of learning rate on both the values simultaneously, we chose learning rate to vary between 0.1 and 0.6 for our training.


Visualizations

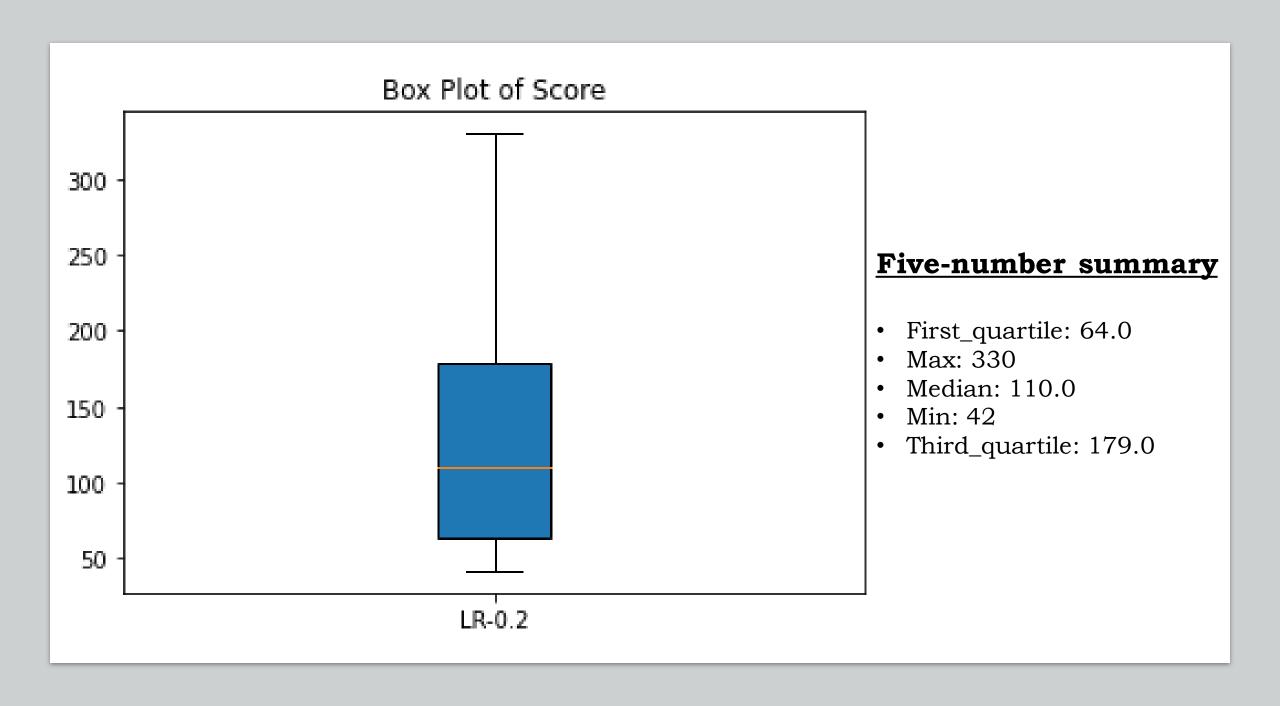
1st Parameter (Learning Rate - 0.1)



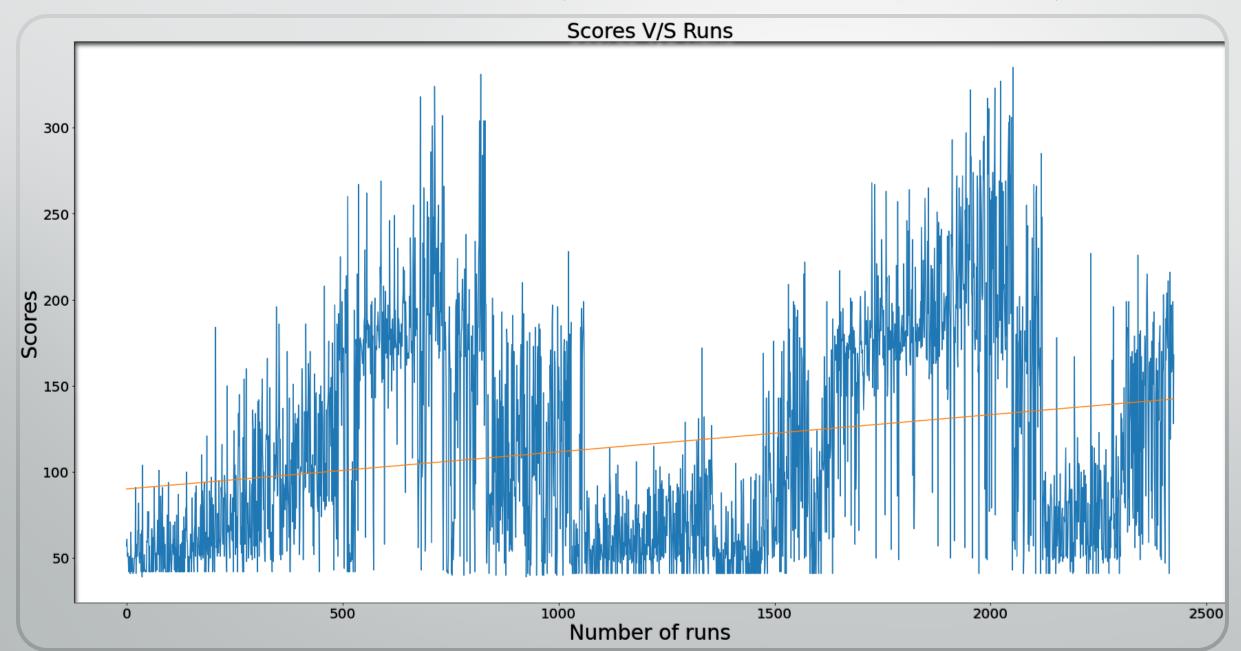


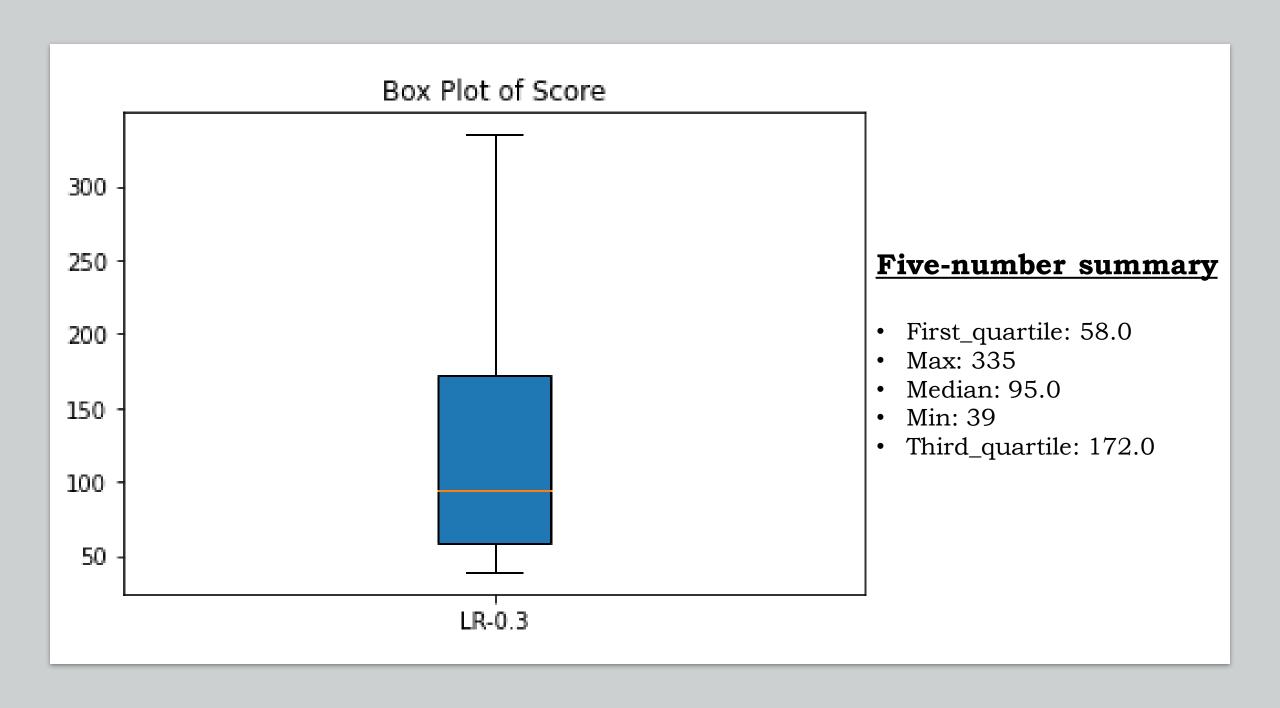
2nd Parameter (Learning Rate - 0.2)







3rd Parameter (Learning Rate - 0.3)



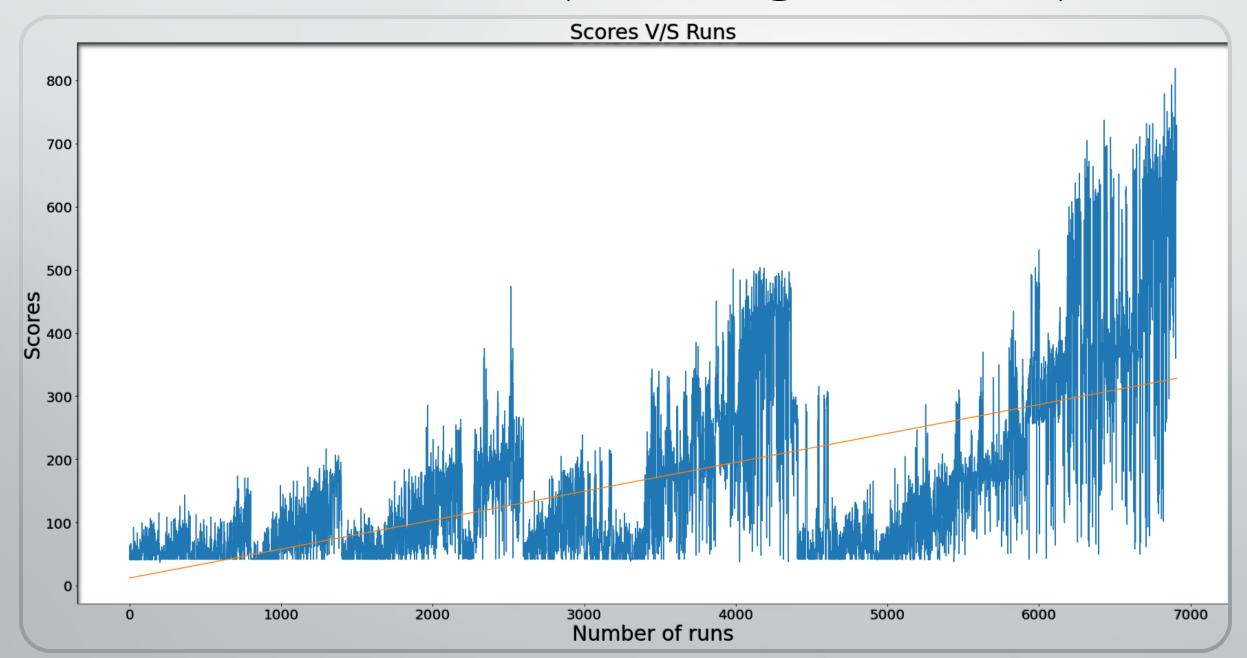



4th Parameter (Learning Rate - 0.4)







5th Parameter (Learning Rate - 0.5)






6th Parameter (Learning Rate - 0.6)





Combined Box Plot



Summarization

Learning Rate	0.1	0.2	0.3	0.4	0.5	0.6
No. of runs	4090	2621	2425	2060	6910	5400
Max Score of all runs	628	330	335	444	819	527
Min Score of all runs	34	42	39	41	37	37
Mean	295.46	127.83	116.27	146.4	170.48	138.89
Median	286.0	110.0	95.0	150.0	113.0	121.0
Standard Deviation	155.29	70.26	64.91	81.97	149.87	87.01

Interpretation

- 0.1 came out to be the best learning rate which was kind of expected, as learning rate 0.1 means that we are giving higher weightage to the current state rather than the future states.
- The problem of not getting all the states in one single run is visible through graphs, as there is a sharp decrease in the score in the subsequent runs.
- Learning for the higher number of runs results in better scores.
- Even though the highest score of learning rate 0.5 is more than that of learning rate 0.1, and that of 0.6 is comparable to that of learning rate 0.1 but that does not mean 0.5 & 0.6 learning rates are better:
 - They are executed for more number of runs.
 - Mean and Median of scores of 0.5 and 0.6 learning rate are much lower than that of 0.1, which means 0.5 and 0.6 are performing good only for a few number of runs, whereas 0.1 is performing good for most number of runs.

Challenges faced during the training phase?

Huge amount of states and the combinations of states present.

Required a lot of time to initialize Q-Table with the new states.

It took around 14 hours to run over 1000 episodes in one go and get most of the states for one particular parameter.

We still cannot be sure that we have seen all the states and stored it in the Q-Table.

References

https://medium.com/acing-ai/how-i-build-an-ai-to-play-dino-run-e37f37bdf153
https://youtu.be/LzaWrmKL1Z4
https://youtu.be/DhdUlDIAG7Y
https://towardsdatascience.com/simple-reinforcement-learning-q-learning-fcddc4b6fe56
https://towardsdatascience.com/a-beginners-guide-to-q-learning-c3e2a3oa653c
https://towardsdatascience.com/math-behind-reinforcement-learning-the-easy-way-1b7edoco3of4
https://towardsdatascience.com/math-of-q-learning-python-code-5dcbdc49b6f6

Thank you